

Customer Segmentation using K-Means Clustering

Indented block

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import datetime as dt
from sklearn.preprocessing import MinMaxScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score, calinski_harabasz_score
```

```
In [30]: from google.colab import drive
drive.mount("/content/gdrive")
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

```
In [31]: import pandas as pd
dataset = pd.read_excel('/content/gdrive/MyDrive/online_retail.xlsx')
```

Data Set Information:

This Online Retail II data set contains all the transactions occurring for a UK-based and registered, non-store online retail between 01/12/2009 and 09/12/2011. The company mainly sells unique all-occasion gift-ware. Many customers of the company are wholesalers.

Attribute Information:

InvoiceNo: Invoice number. Nominal. A 6-digit integral number uniquely assigned to each transaction. If this code starts with the letter 'c', it indicates a cancellation.

StockCode: Product (item) code. Nominal. A 5-digit integral number uniquely assigned to each distinct product.

Description: Product (item) name. Nominal.

Quantity: The quantities of each product (item) per transaction. Numeric.

InvoiceDate: Invoice date and time. Numeric. The day and time when a transaction was generated.

UnitPrice: Unit price. Numeric. Product price per unit in sterling (£). (1 Pound sterling equals 96.22 Indian Rupee)

CustomerID: Customer number. Nominal. A 5-digit integral number uniquely assigned to each customer.

Country: Country name. Nominal. The name of the country where a customer resides.

```
In [ ]: dataset.columns
```

```
Out[ ]: Index(['Invoice', 'StockCode', 'Description', 'Quantity', 'InvoiceDate',
          'Price', 'Customer ID', 'Country', 'month_year'],
```

dtype='object')

```
In [ ]: dataset.shape
```

Out[]: (525461, 8)

```
In [ ]: dataset.head(10)
```

Out[]:

| | Invoice | StockCode | Description | Quantity | InvoiceDate | Price | Customer ID | Country |
|---|---------|-----------|-------------------------------------|----------|---------------------|-------|-------------|----------------|
| 0 | 489434 | 85048 | 15CM CHRISTMAS GLASS BALL 20 LIGHTS | 12 | 2009-12-01 07:45:00 | 6.95 | 13085.0 | United Kingdom |
| 1 | 489434 | 79323P | PINK CHERRY LIGHTS | 12 | 2009-12-01 07:45:00 | 6.75 | 13085.0 | United Kingdom |
| 2 | 489434 | 79323W | WHITE CHERRY LIGHTS | 12 | 2009-12-01 07:45:00 | 6.75 | 13085.0 | United Kingdom |
| 3 | 489434 | 22041 | RECORD FRAME 7" SINGLE SIZE | 48 | 2009-12-01 07:45:00 | 2.10 | 13085.0 | United Kingdom |
| 4 | 489434 | 21232 | STRAWBERRY CERAMIC TRINKET BOX | 24 | 2009-12-01 07:45:00 | 1.25 | 13085.0 | United Kingdom |
| 5 | 489434 | 22064 | PINK DOUGHNUT TRINKET POT | 24 | 2009-12-01 07:45:00 | 1.65 | 13085.0 | United Kingdom |
| 6 | 489434 | 21871 | SAVE THE PLANET MUG | 24 | 2009-12-01 07:45:00 | 1.25 | 13085.0 | United Kingdom |
| 7 | 489434 | 21523 | FANCY FONT HOME SWEET HOME DOORMAT | 10 | 2009-12-01 07:45:00 | 5.95 | 13085.0 | United Kingdom |
| 8 | 489435 | 22350 | CAT BOWL | 12 | 2009-12-01 07:46:00 | 2.55 | 13085.0 | United Kingdom |
| 9 | 489435 | 22349 | DOG BOWL , CHASING BALL DESIGN | 12 | 2009-12-01 07:46:00 | 3.75 | 13085.0 | United Kingdom |

```
In [10]: print("Number of customers are ",len(dataset["Customer ID"].value_counts()))
```

Number of customers are 4314

```
In [32]: print(dataset.isnull().sum())
dataset = dataset.dropna()
print(dataset.isnull().sum())
print(dataset.shape)
```

```
Invoice      0
StockCode    0
Description  2928
Quantity     0
InvoiceDate  0
Price        0
Customer ID  107927
Country      0
dtype: int64
Invoice      0
StockCode    0
```

Description 0
Quantity 0
InvoiceDate 0
Price 0
Customer ID 0
Country 0
dtype: int64
(417534, 8)

```
In [33]: dataset = dataset[(dataset['Quantity']>0)]
```

```
In [34]: dataset.describe()
```

Out[34]:

| | Quantity | Price | Customer ID |
|-------|---------------|---------------|---------------|
| count | 407695.000000 | 407695.000000 | 407695.000000 |
| mean | 13.586686 | 3.294188 | 15368.504107 |
| std | 96.842229 | 34.756655 | 1679.795700 |
| min | 1.000000 | 0.000000 | 12346.000000 |
| 25% | 2.000000 | 1.250000 | 13997.000000 |
| 50% | 5.000000 | 1.950000 | 15321.000000 |
| 75% | 12.000000 | 3.750000 | 16812.000000 |
| max | 19152.000000 | 10953.500000 | 18287.000000 |

```
In [35]: Country_quantity = dataset["Country"].value_counts()  
Country_quantity
```

Out[35]:


| | |
|----------------------|--------|
| United Kingdom | 370951 |
| EIRE | 8507 |
| Germany | 7661 |
| France | 5470 |
| Netherlands | 2730 |
| Spain | 1235 |
| Switzerland | 1170 |
| Belgium | 1038 |
| Portugal | 984 |
| Sweden | 868 |
| Channel Islands | 821 |
| Italy | 710 |
| Australia | 630 |
| Cyprus | 541 |
| Austria | 524 |
| Greece | 512 |
| Denmark | 418 |
| Norway | 365 |
| Finland | 347 |
| United Arab Emirates | 315 |
| Unspecified | 277 |
| USA | 230 |
| Poland | 182 |
| Malta | 170 |
| Japan | 164 |
| Lithuania | 154 |
| Singapore | 117 |
| Canada | 77 |

| | |
|-------------|----|
| Thailand | 76 |
| Israel | 74 |
| Iceland | 71 |
| RSA | 65 |
| Brazil | 62 |
| West Indies | 54 |
| Korea | 53 |
| Bahrain | 42 |
| Nigeria | 30 |

Name: Country, dtype: int64

In [36]:

```
import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize = (10,7))
sns.set_theme(style="whitegrid")
plt.xticks(rotation=80,size=13)
sns.barplot(x=Country_quantity.index[1:], y=Country_quantity[1:])
plt.show()
```




In [37]:

```
import datetime as dt
# month_info = pd.DataFrame(dataset['InvoiceDate'].dt.month)
# year_info = pd.DataFrame(dataset['InvoiceDate'].dt.year)
# new = pd.concat([month_info,year_info],axis=1)
#temp = pd.to_datetime(month_info,year_info.assign(Day=1))
#temp= pd.to_datetime(dataset[['year', 'month']].assign(Day=1))
# dataset['revenue'] = dataset['Price'] * dataset['Quantity']
temp = list(map(lambda x: x.strftime("%Y-%m"),dataset["InvoiceDate"].copy()))
#temp = pd.to_datetime(dataset["InvoiceDate"].copy(),format="%Y-%m-%d")
dataset["month_year"] = temp
```

```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns


plot = pd.DataFrame(dataset.groupby(['month_year'])['Invoice'].count()).reset_index()
plt.figure(figsize=(10,5))
ax = sns.lineplot(x="month_year", y="Invoice", data = plot)
plt.show()
```




```
In [23]: dataset["revenue"] = dataset["Price"]*dataset["Quantity"]
dataset["revenue"]
```

```
Out[23]: 0      83.40
1      81.00
2      81.00
3     100.80
4      30.00
...
525456    5.90
525457    3.75
525458    3.75
525459    7.50
525460    3.90
Name: revenue, Length: 407695, dtype: float64
```

```
In [ ]: data_temp = pd.DataFrame(dataset.groupby(['month_year'])['revenue'].sum()).reset_index()
plt.figure(figsize=(10,5))
ax = sns.lineplot(x = 'month_year', y='revenue', data = data_temp)
```



```
In [ ]: data_temp2 = pd.DataFrame(dataset.groupby(['Country'])['revenue'].sum()).reset_index
plt.figure(figsize=(15,5))
ax=sns.barplot(x='Country', y='revenue',data=data_temp2)
plt.xticks(rotation=65,size=10)
plt.show()
```



```
In [ ]: dataset.columns
```

```
Out [ ]: Index(['Invoice', 'StockCode', 'Description', 'Quantity', 'InvoiceDate',
        'Price', 'Customer ID', 'Country', 'month', 'year', 'month_year',
        'revenue'],
        dtype='object')
```

RFM Analysis:

- R (Recency) - most recent purchase
- F (Frequency) - how frequent is the purchase
- M (Monetary) - amount spent in each purchase

```
In [26]:
```

```
dataset['Customer ID'] = dataset['Customer ID'].astype(str)
dataset['Amount'] = dataset['Quantity']*dataset['Price']
rfm_dataset_monetary = dataset.groupby('Customer ID')['Amount'].sum()
rfm_dataset_monetary.reset_index()
rfm_dataset_monetary.columns = ['Customer ID', 'Amount']
print(rfm_dataset_monetary)
```

```
Customer ID
12346.0      372.86
12347.0     1323.32
12348.0      222.16
12349.0     2671.14
12351.0      300.93
...
18283.0      641.77
18284.0      461.68
18285.0      427.00
18286.0     1296.43
18287.0     2345.71
Name: Amount, Length: 4314, dtype: float64
```

In []:

```
rfm_dataset_frequency = dataset[['Customer ID', 'Invoice']].groupby(['Customer ID', '
rfm_dataset_frequency = rfm_dataset_frequency.reset_index().groupby(["Customer ID"])
rfm_dataset_frequency.columns = ['Customer ID', 'Frequency']
print(rfm_dataset_frequency)
```

```
Customer ID  Frequency
0      12346.0         11
1      12347.0          2
2      12348.0          1
3      12349.0          3
4      12351.0          1
...      ...      ...
4309     18283.0          6
4310     18284.0          1
4311     18285.0          1
4312     18286.0          2
4313     18287.0          4
```

[4314 rows x 2 columns]

In [38]:

```
dataset['InvoiceDate'] = pd.to_datetime(dataset['InvoiceDate'], format='%d-%m-%Y %H:%
max_date = max(dataset['InvoiceDate'])
dataset['Diff'] = max_date - dataset['InvoiceDate']
rfm_dataset_recency = dataset.groupby('Customer ID')['Diff'].min()
rfm_dataset_recency = rfm_dataset_recency.reset_index()
rfm_dataset_recency.columns = ['Customer ID', 'Diff']
rfm_dataset_recency['Diff'] = rfm_dataset_recency['Diff'].dt.days
print(rfm_dataset_recency)
```

```
Customer ID  Diff
0      12346.0   164
1      12347.0     2
2      12348.0    73
3      12349.0    42
4      12351.0    10
...      ...      ...
4309     18283.0    17
4310     18284.0    66
4311     18285.0   295
4312     18286.0   111
4313     18287.0    17
```

[4314 rows x 2 columns]

```
In [ ]: rfm_dataset_final = pd.merge(rfm_dataset_monetary, rfm_dataset_frequency, on='Customer ID')
rfm_dataset_final = pd.merge(rfm_dataset_final, rfm_dataset_recency, on='Customer ID')
rfm_dataset_final.columns = ['Customer ID', 'Amount', 'Frequency', 'Recency']
print(rfm_dataset_final)
```

| | Customer ID | Amount | Frequency | Recency |
|------|-------------|---------|-----------|---------|
| 0 | 12346.0 | 372.86 | 11 | 164 |
| 1 | 12347.0 | 1323.32 | 2 | 2 |
| 2 | 12348.0 | 222.16 | 1 | 73 |
| 3 | 12349.0 | 2671.14 | 3 | 42 |
| 4 | 12351.0 | 300.93 | 1 | 10 |
| ... | ... | ... | ... | ... |
| 4309 | 18283.0 | 641.77 | 6 | 17 |
| 4310 | 18284.0 | 461.68 | 1 | 66 |
| 4311 | 18285.0 | 427.00 | 1 | 295 |
| 4312 | 18286.0 | 1296.43 | 2 | 111 |
| 4313 | 18287.0 | 2345.71 | 4 | 17 |

[4314 rows x 4 columns]

```
In [ ]: rfm_dataset_final.shape
```

```
Out[ ]: (4314, 4)
```

```
In [ ]: Q1 = rfm_dataset_final.Amount.quantile(0.05)
Q3 = rfm_dataset_final.Amount.quantile(0.95)
IQR = Q3 - Q1
rfm_dataset_final = rfm_dataset_final[(rfm_dataset_final.Amount >= Q1 - 1.5*IQR) & (rfm_dataset_final.Amount <= Q3 + 1.5*IQR)]

Q1 = rfm_dataset_final.Recency.quantile(0.05)
Q3 = rfm_dataset_final.Recency.quantile(0.95)
IQR = Q3 - Q1
rfm_dataset_final = rfm_dataset_final[(rfm_dataset_final.Recency >= Q1 - 1.5*IQR) & (rfm_dataset_final.Recency <= Q3 + 1.5*IQR)]

Q1 = rfm_dataset_final.Frequency.quantile(0.05)
Q3 = rfm_dataset_final.Frequency.quantile(0.95)
IQR = Q3 - Q1
rfm_dataset_final = rfm_dataset_final[(rfm_dataset_final.Frequency >= Q1 - 1.5*IQR) & (rfm_dataset_final.Frequency <= Q3 + 1.5*IQR)]
```

```
In [ ]: X = rfm_dataset_final[['Amount', 'Frequency', 'Recency']]
scaler = MinMaxScaler()
rfm_dataset_scaled = scaler.fit_transform(X)
```

```
In [ ]: rfm_dataset_scaled = pd.DataFrame(rfm_dataset_scaled)
rfm_dataset_scaled.columns = ['Amount', 'Frequency', 'Recency']
rfm_dataset_scaled.head()
```


```
Out[ ]:
```

| | Amount | Frequency | Recency |
|---|----------|-----------|----------|
| 0 | 0.024629 | 0.370370 | 0.439678 |
| 1 | 0.087409 | 0.037037 | 0.005362 |
| 2 | 0.014674 | 0.000000 | 0.195710 |
| 3 | 0.176437 | 0.074074 | 0.112601 |

| | Amount | Frequency | Recency |
|---|----------|-----------|----------|
| 4 | 0.019877 | 0.000000 | 0.026810 |

In []:

```
g = sns.PairGrid(rfm_dataset_scaled)
g.map(sns.scatterplot);
```



Segmentation based on Amount, Recency and Frequency

In []:

```
within_sum_square = []
range_n_clusters = [i for i in range(2,11)]
for num_clusters in range_n_clusters:
    kmeans = KMeans(n_clusters=num_clusters, init='k-means++', max_iter=300, random_
    within_sum_square.append(kmeans.inertia_)
    cluster_labels = kmeans.labels_
    silhouette_avg = silhouette_score(rfm_dataset_scaled, cluster_labels)
    # c_avg = calinski_harabasz_score(rfm_dataset_scaled, cluster_labels)
    print("For n_clusters={0}, the silhouette score is {1}".format(num_clusters, sil
```


```
For n_clusters=2, the silhouette score is 0.5661945492047095
For n_clusters=3, the silhouette score is 0.5477575520199691
For n_clusters=4, the silhouette score is 0.4743336103129437
For n_clusters=5, the silhouette score is 0.4377793045333027
For n_clusters=6, the silhouette score is 0.3880489018428877
For n_clusters=7, the silhouette score is 0.381351139040909
```

For n_clusters=8, the silhouette score is 0.387409124390399
 For n_clusters=9, the silhouette score is 0.37027024047295526
 For n_clusters=10, the silhouette score is 0.35262500025185517

```
In [ ]: kmeans = KMeans(n_clusters=3, max_iter=300,init="k-means++",random_state=42)
kmeans.fit(rfm_dataset_scaled)
lbs = kmeans.labels_
print(kmeans.labels_)
```

```
[1 0 0 ... 1 0 0]
```

```
In [ ]: plt.plot(range_n_clusters,within_sum_square)
plt.show()
```



```
In [ ]: rfm_dataset_final['Cluster_Id'] = lbs
rfm_dataset_final.head()
```


```
Out[ ]:
```

| | Customer ID | Amount | Frequency | Recency | Cluster_Id |
|---|-------------|---------|-----------|---------|------------|
| 0 | 12346.0 | 372.86 | 11 | 164 | 1 |
| 1 | 12347.0 | 1323.32 | 2 | 2 | 0 |
| 2 | 12348.0 | 222.16 | 1 | 73 | 0 |
| 3 | 12349.0 | 2671.14 | 3 | 42 | 0 |
| 4 | 12351.0 | 300.93 | 1 | 10 | 0 |

```
In [ ]: from mpl_toolkits.mplot3d import Axes3D
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

fig = plt.figure(figsize=(20,10))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(rfm_dataset_final["Amount"][rfm_dataset_final.Cluster_Id == 0], rfm_dataset_final["Frequency"][rfm_dataset_final.Cluster_Id == 0], rfm_dataset_final["Recency"][rfm_dataset_final.Cluster_Id == 0], rfm_dataset_final["Cluster_Id"][rfm_dataset_final.Cluster_Id == 0])
ax.scatter(rfm_dataset_final["Amount"][rfm_dataset_final.Cluster_Id == 1], rfm_dataset_final["Frequency"][rfm_dataset_final.Cluster_Id == 1], rfm_dataset_final["Recency"][rfm_dataset_final.Cluster_Id == 1], rfm_dataset_final["Cluster_Id"][rfm_dataset_final.Cluster_Id == 1])
ax.scatter(rfm_dataset_final["Amount"][rfm_dataset_final.Cluster_Id == 2], rfm_dataset_final["Frequency"][rfm_dataset_final.Cluster_Id == 2], rfm_dataset_final["Recency"][rfm_dataset_final.Cluster_Id == 2], rfm_dataset_final["Cluster_Id"][rfm_dataset_final.Cluster_Id == 2])
ax.view_init(30, 185)
plt.legend(("Cluster 0 ", "Cluster 1", "Cluster 2"))
plt.xlabel("Amount")
plt.ylabel("Frequency")
```

```
ax.set_zlabel('Recency')
plt.show()
```



Segmentation based on Amount and Frequency

```
In [ ]: rfm_dataset_af = rfm_dataset_scaled[["Amount", "Frequency"]].copy()
rfm_dataset_af
```

```
Out [ ]:
```

| | Amount | Frequency |
|------|----------|-----------|
| 0 | 0.024629 | 0.370370 |
| 1 | 0.087409 | 0.037037 |
| 2 | 0.014674 | 0.000000 |
| 3 | 0.176437 | 0.074074 |
| 4 | 0.019877 | 0.000000 |
| ... | ... | ... |
| 4232 | 0.042391 | 0.185185 |
| 4233 | 0.030495 | 0.000000 |
| 4234 | 0.028205 | 0.000000 |
| 4235 | 0.085633 | 0.037037 |
| 4236 | 0.154941 | 0.111111 |

4237 rows × 2 columns


```
In [ ]: d2kmeans = KMeans(n_clusters = 3,          # Set amount of clusters
                          init = 'k-means++',      # Initialization method for kmeans
                          max_iter = 100)         # Maximum number of iterations

pred_y = d2kmeans.fit_predict(rfm_dataset_af)

# Plot the data
plt.scatter(rfm_dataset_scaled["Amount"],
```

```
rfm_dataset_scaled["Frequency"])

# Plot the clusters
plt.scatter(d2kmeans.cluster_centers[:, 0],
            d2kmeans.cluster_centers[:, 1],
            s=200,                                # Set centroid size
            c='red')
plt.legend(["data", "centroid"])
plt.xlabel("Amount")
plt.ylabel("Frequency")                          # Set centroid color
plt.show()
```




```
In [ ]: #pred_y
rfm_copy = rfm_dataset_af.copy()
rfm_copy["Cluster_Id"] = pred_y
rfm_copy.head()
```

```
Out[ ]:   Amount  Frequency  Cluster_Id
0  0.024629   0.370370         1
1  0.087409   0.037037         0
2  0.014674   0.000000         0
3  0.176437   0.074074         0
4  0.019877   0.000000         0
```

```
In [ ]: u_labels = np.unique(pred_y)


#plotting the results:

for i in u_labels:
    plt.scatter(rfm_copy[pred_y == i]['Amount'], rfm_copy[pred_y == i]['Frequency'], label=i)
    plt.scatter(d2kmeans.cluster_centers[:, 0],
                d2kmeans.cluster_centers[:, 1],
                s=100,                                # Set centroid size
                c='red')
plt.legend()
plt.xlabel("Amount")
plt.ylabel("Frequency")
plt.show()
```




Plots that summarize the Output

```
In [ ]: sns.boxplot(x='Cluster_Id', y='Amount', data=rfm_dataset_final)
plt.show()
```



```
In [ ]: sns.boxplot(x='Cluster_Id', y='Frequency', data=rfm_dataset_final)
plt.show()
```



```
In [ ]: sns.boxplot(x='Cluster_Id', y='Recency', data=rfm_dataset_final)  
plt.show()
```

